

# Biologically-Inspired Computation (F20BC/F21BC)

## Coursework

This assessment is worth 40% of your overall course mark

**Due: 3:30pm local time, Friday 22<sup>nd</sup> November 2024**

### Overview

This assessment aims to increase your understanding of both artificial neural networks (ANNs) and particle swarm optimisation (PSO), two biologically-inspired techniques which are taught in the course. It involves implementing both ANNs and PSO from scratch, and experimentally investigating how PSO can be used to optimise an ANN to carry out a specified task.

Please read through the following important points before you begin:

- In order to encourage discussion, **the assessment involves working in pairs**, i.e. 2 people. You need to choose a partner from the same level (i.e. an F20BC student cannot be paired with an F21BC student). You should report your team choice on Canvas by **Friday 27<sup>th</sup> September**:
  - For Edinburgh students, see Edinburgh Information → Edinburgh Team Formation
  - For Dubai students, see Dubai Information → Dubai Team Formation
- Each member of a pair should contribute equally. We will be asking you to summarise your contributions when you submit your work, and we will rebalance the marks within a pair in cases where one member contributes substantially more than the other.
- **Do not copy code from the internet or from other students.** This is very easy for us to detect, and every year we have to report students to the academic misconduct board for plagiarism. This involves a lot of work for us, and it is very stressful for students to go through this process, so please *please* don't submit code that isn't your own without acknowledging its source!

**This is assessed coursework.** You are allowed to discuss this assignment with students outside of your pair, but you should not copy their work, and you should not share your own work with other students. We will be carrying out automated plagiarism checks on both code and text submissions.

**Special note for using existing code or generative AI.** If you are using code that you have not yourself written, then this must clearly be indicated, making clear which parts were not written by you and clearly stating where it came from. If your code is found elsewhere by the person marking your work (including in the output of large language models) and you have not mentioned this, you may find yourself having to go before a disciplinary committee.

**Late submission and extensions.** Late submissions will be marked according to the university's late submissions policy, i.e. a 30% deduction if submitted within 5 working days of the deadline, and a mark of 0% after that. The deadline for this work is not negotiable. If you are unable to complete the assignment by the deadline due to circumstances beyond your control (e.g. illness or family bereavement), you should complete and submit a mitigating circumstances application: <https://www.hw.ac.uk/students/studies/examinations/mitigating-circumstances.htm>

What you are asked to do:

1. Implement a multi-layer ANN architecture
2. Implement a PSO algorithm
3. Use PSO to train the parameters of the ANN
4. Train the ANN to solve a specified problem
5. Experimentally investigate the effect of ANN and PSO hyperparameters
6. Write a short report and submit both the report and your code to Canvas
7. Indicate your individual contributions on the group signing sheet

These tasks are described in more detail below. Implementation should be done using a language of your choice (e.g. Java, Python, Matlab, C, C++). The aim is for you to learn how to implement biologically-inspired approaches from scratch, so you are strongly encouraged not to use existing ANN or PSO libraries, since doing so will considerably reduce your marks for implementation. Whilst you may consult web pages and large language models for guidance, the code you submit should be your own, and you will be penalised if it strongly resembles code found on the web or generated by large language models. Note that we use various tools to check for this.

**You should start work on task 1 of the assessment once ANNs have been covered in lectures. At that point, you can also write the evaluation code for task 4, and can test this code on ANNs with random weights. PSO will not be covered in lectures until week 7, though if you are eager to get started on this part, the introductory lecture on PSO is available under Week 7 on Canvas.**

### 1. Implement a multi-layer ANN architecture

You should implement a simple feedforward multilayer architecture. The number of nodes in each layer, the number of layers, and the activation function used in each layer should be configurable.

Note that your ANN will be applied to a regression task in task 4, and that the number of nodes in the input and output layers will be determined by this.

You do not need to implement any classical training algorithms, e.g. backpropagation, since you will be using PSO to train the parameters (i.e. weights and biases) of the model.

Here is a list of activation functions that should be implemented, though you may also investigate other suitable functions:

- Logistic function:  $\frac{1}{1+e^{-x}}$
- ReLU (rectified linear unit):  $\max(0, x)$
- Hyperbolic tangent:  $\tanh x$

### 2. Implement a PSO algorithm

You should implement Particle Swarm Optimisation (PSO). Specifically, you should implement the form of PSO described in Section 3.5 of the book “Essentials of Metaheuristics”, which is also described in the Week 7 PSO lecture: <https://cs.gmu.edu/~sean/book/metaheuristics/>

Unlike the original version of PSO (which you may find described on websites or by large language models), this form of PSO uses informants. That is, each particle should be influenced by a subset of the other particles in the swarm, rather than just the swarm best. See the lecture for more information on this.

You are also welcome to extend these minimum requirements in some way. For example, you might implement boundary handling or consider other approaches you come across in the PSO literature.

### 3. Use PSO to train the parameters of the ANN

PSO is an optimisation algorithm that optimises a vector of values. In this case, this vector of values should represent the set of parameters (i.e. weights and biases) of an ANN, and the goal is to find values for each of these parameters so that the ANN correctly solves a specific problem (see task 4 below for details about this). So, you need to use the PSO code you wrote in task 2 to optimise the parameters of the ANN code you wrote in task 1, i.e. you'll need to figure out how to couple together your PSO and ANN code. There will be some hints on how to do this in the in-person lecture in Week 7.

Here are some points to bear in mind:

- Each particle within the PSO swarm represents an ANN as a fixed-length vector of floating-point values, each of which encodes the value of a particular parameter of the ANN.
- Each time a particle is evaluated in PSO, the values in its vector should be used to set the parameters of the ANN, and the ANN should then be evaluated in order to measure its accuracy on the given problem. This accuracy value then becomes the particle's fitness.
- Although there are versions of PSO that can handle variable-length vectors, you are not expected to know about these. Consequently, the architecture of the ANN (i.e. the number of layers and neurons) should be specified at the beginning of a PSO run and remain fixed.
- One advantage of using PSO, rather than backpropagation, is that you can optimise other aspects of the ANN in addition to the weights and biases. For instance, you are encouraged to also try encoding the activation functions used by the ANN within the PSO solution vector, and extra marks will be available if you do this.

### 4. Train the ANN to solve a specified problem

The problem domain for this work is regression, i.e. training an ANN to predict the correct numeric output for a given set of input values. You will be using this dataset:

<https://archive.ics.uci.edu/dataset/165/concrete+compressive+strength>

Basically, the aim is to predict the compressive strength of a sample of concrete given 8 numeric features that describe its characteristics. That is, your ANN will take 8 numeric values as inputs, and its output will be its prediction of the compressive strength.

You can download a CSV (comma-separated values) version of the data set from the following link; each row represents one sample, with the first 8 numbers being the input features and the last number being the output value that you want your ANN to predict; there's also a header row:

<https://www.kaggle.com/datasets/elikplim/concrete-compressive-strength-data-set>

The dataset comprises 1030 instances/samples (i.e. rows). You should split this into a training set and a test set, using the former to train your neural network, and the latter to measure how well it generalises. A typical split is 70% of the data used for training, and 30% used for testing. You're welcome to use more advanced forms of evaluation if you like, such as cross-validation, but are not expected to do so. If you've not done any data science before, there are plenty of tutorials online that describe the basics. For example, you might find these links useful:

<https://developers.google.com/machine-learning/crash-course/overfitting/dividing-datasets>

<https://developers.google.com/machine-learning/crash-course/overfitting/overfitting>

To evaluate how well your ANN works, you should use a regression metric. The simplest is the mean absolute error (MAE), which is basically the average error between the expected outputs (i.e. the last column of the CSV file) and the outputs your trained ANN actually predicts on the test set. Again, there is plenty of information online, including:

<https://towardsdatascience.com/what-are-rmse-and-mae-e405ce230383>

## 5. Experimentally investigate the effect of ANN and PSO hyperparameters

As discussed in the lectures, both ANNs and PSO have numerous hyperparameters. For ANNs, this includes the number of layers, the number of neurons in a layer, and the activation function used in each layer. For PSO, it includes the values of the acceleration coefficients, the number of informants, the swarm size and the number of iterations. The aim of this experimental investigation is to get some insight into how these values effect the ability of PSO to optimise an ANN than correctly solves the specified problem.

Specifically, you are asked to try and answer the following questions:

- What effect does the ANN architecture have on its ability to solve the problem? Architecture refers to aspects such as the number of layers, the number of neurons in a layer, and the activation function used in each layer.
- What is the best way of allocating solution evaluations? That is, the number of solution evaluations in a run of PSO is the product of the swarm size and the number of iterations. For a fixed budget of evaluations, say 500, would it be better to have a swarm of size 10 with 50 iterations, a swarm of size 50 with 10 iterations, or somewhere in between?
- What is the effect of varying the acceleration coefficients in the PSO velocity update equation? That is, look at how different balances between the coefficients affect PSO's ability to solve this particular problem. Of particular interest is the balance between the social and cognitive components, but you might also look at the other coefficients if you have time. Looking at the PSO literature might help you to find sensible values for these.

Here are some points to bear in mind:

- Pick a sensible range of values for each hyperparameter that you investigate. This could be guided by values you find in lectures, books, published papers etc. You can also use a more systematic approach, such as random search or grid search, though this is not required.
- PSO is a stochastic algorithm. That is, each time you run it, the particles start in different positions and how they move has a random element. This means that, for the same set of hyperparameter values, you will likely get different results each time you run it. To address this, it is strongly recommended that you carry out at least 10 runs for each set of hyperparameter values, and present the average and standard deviation across these 10 runs, rather than the results of individual runs.

## 6. Write a short report and submit both the report and your code to Canvas

Your report should:

- Be no more than **6 pages** in length (max of 3000 words), excluding references and appendices. You should take this into account when planning your experiments. If you have more results than you have space for, then pick the results that you think are most insightful and briefly mention which other experiments you carried out.
- Be written in Arial, or a similar font, with a minimum **font size of 12**.

It should contain the following sections:

- **Implementation:** Briefly describe your code. For instance, what does each function do? Describe the motivation behind any design decisions you made, and note any interesting aspects of your implementation. You can assume the reader already knows the basics of ANNs and PSO, so there's no need to introduce these in your report.
- **Experiments and Results:** State which questions you attempted to answer, which hyperparameters you investigated, and indicate which value ranges you looked at and why you chose these ranges. You can use references to motivate this. Then, using tables and plots, show your results. When carrying out multiple runs to compensate for stochasticity (see above), tables should show averages (and ideally standard deviations) rather than results of individual runs. Plots are useful for illustrating trends that can be hard to spot in tables alone.
- **Discussion and Conclusions:** What did you discover? How well did your ANNs solve the problem, and how did the hyperparameter settings affect this? Don't just describe what you saw; also try to explain why this is the case, e.g. why do you think a particular hyperparameter had a particular effect? You can use references to the literature to support your arguments.
- **References:** These should be in a standard format, e.g. Harvard style, and should be cited in the earlier sections. Don't include references that you don't use.
- **Appendices (optional):** You can include additional results here, if you want to. However, your marker won't look at this section in any detail, so anything important and key to your discussion should appear in the earlier sections.

You should submit both your report (as a **pdf** file) and your code (as a **zip** file) to Canvas using the links provided. Make sure it's clear how to run your code.

## 7. Indicate your individual contributions on the group signing sheet

The Coursework Group Signing Sheet is available on Canvas. Both members of your team should outline their contribution to the coursework, indicating which parts of the code, investigation and report they contributed to, and sign the sheet. It should then either be embedded in your report, or included in the code submission zip file. **No marks will be issued until this has been submitted.**

## Marking and Feedback

See the table on the next page for the marking scheme. Note that the weightings are slightly different for F20BC and F21BC, and reflect a higher expectation in terms of interpretation and analysis for the latter. The grade requirements outlined in the table are indicative only, and we'll take into account your achievements across the assessment when deciding a mark. This means, for instance, if you choose to put extra effort into the implementation part, then we wouldn't expect quite as much work in the experimental part. Conversely, if your code is quite basic, then you can compensate for this by putting more work into the experimental investigation and by making greater use of the literature in your discussion.

Once your work is marked, we will provide you with marks and feedback, and will aim to do this before the end of the winter break. Whilst we will try to do this as fast as possible, please bear in mind that the coursework is due towards the end of the semester, so we will be marking this at the same time as your exam scripts, as well as trying to take some annual leave.

**You can also get informal feedback from lab helpers before you submit your work. In Edinburgh, lab helpers will be available during the Friday lab session for this purpose. They will also be able to help you if you get stuck, so please make use of this resource.**

Criteria	Weight	A (70-100%)	B (60-69%)	C (50-59%)	D (40-49%)	E/F (<40%)
<b>Implementation</b> (i.e. code, comments, description and motivation in the report)	50% for F20BC  45% for F21BC	Creative implementations of ANN and PSO that exceed the basic requirements. ANN and PSO are correctly coupled together. Correct evaluation code. Easy to read and well structured.	Correct implementations of the requirements. ANN and PSO are correctly coupled together. Generally good coding, structure and documentation.	Some significant issues in terms of correctness, structure, coding practice (including lack of comments) and documentation. ANN and PSO are correctly coupled together.	Major issues in terms of correctness, structure, coding practice and documentation and/or ANN and PSO are not correctly coupled together.	Critical errors: for example, the code does not compile and/or run, or inappropriate algorithms have been implemented.
<b>Experimental study</b> (i.e. choice and validity of experiments performed, presentation of results)	20%	Suitable hyperparameters are investigated and their values are well chosen. Suitable results have been collected and are informatively presented and meaningful.	Some minor issues in terms of hyperparameters investigated and their values, the experiments performed, or the presentation of results.	Significant issues in terms of hyperparameters investigated and their values, the experiments performed, or the presentation of results.	Some major issues: experiments do not make sense, have invalid results, or the study is not adequately described.	Some critical issues: experimental study is nonsensical or missing, inappropriate experiments, or the description of the study is uninformative.
<b>Discussion</b> (i.e. results interpretation, understanding demonstrated, conclusions, use of the wider literature or internet resources)	20% for F20BC  25% for F21BC	Clear, insightful discussion that shows a good understanding of ANNs and PSO and includes well chosen references to the wider literature or internet resources.	Generally clear and insightful, but shows some misunderstanding. Adequate use of the wider literature or internet resources.	Discussion is limited in terms of the depth or volume of understanding it demonstrates. Little or no use of the wider literature or internet resources.	Some major issues in terms of depth or volume of understanding. No use of the wider literature or internet resources.	No real demonstration that the subject matter has been understood, or very limited in its scope.
<b>Report</b> (i.e. structure, language, referencing etc.)	10%	Report is well structured and correctly divided into sections; readable, with good use of language; respects page limit and formatting guidelines	Report is suitably structured and divided into sections; mostly readable with good use of language; respects page limit and formatting guidelines	Report is structured but not divided into sections; language issues that affect readability; notable formatting issues	Report is poorly structured; poor readability; significant formatting issues	Report has a nonsensical structure; very hard to read; problematic formatting